


RSI pilot learning curves | development seed 17

- Train segmentation composite (augmented view)


- - - Train total objective (segmentation + weighted risk)
- Validation Dice (original image coordinates)

★ Best epoch selected by validation Dice


A: CNN
DONE 100/100




A: Transformer
DONE 100/100




B: interaction
DONE 40/40




D0
DONE 40/40




RSI ($\lambda = 1$)
DONE 40/40




Mean-Hinge ($\lambda = 1$)
DONE 40/40



Abs-Hard ($\lambda = 1$)
DONE 40/40



Measured stage wall time
DONE.json only



Composite = 0.5 BCE + 0.5 (1 - SoftDice), letterbox valid region. Training curves use augmented views; no smoothing.

Validation: original-image hard Dice, reference mean within image, then image mean; used for development checkpoint selection.

Panels use independent y-axis ranges. Stage wall times are recorded elapsed times, not GPU utilization or a uniform GPU-hour budget. Test scoring remains locked.